Lavkush Kumar

Hyderabad, India | +91-9122036484 | lav84095@gmail.com | linkedin.com/in/lavkush-kumar | github.com/lavkush-kumar

PROFESSIONAL SUMMARY

Azure Data Engineer with 2.5 YOE specializing in Azure Databricks, PySpark, Delta Lake, and ADF within enterprise Medallion Lakehouses. Engineered scalable batch and real-time streaming architectures (Apache Kafka, ADLS Gen2, Unity Catalog) processing 2M+ events/day into Bronze, Silver, and Gold Delta tables with sub-second latency. Automated DataOps CI/CD release cycles and infrastructure provisioning via Terraform, cutting ETL execution runtime by 35% and sustaining 99.95% cloud data platform availability.

TECHNICAL SKILLS

Data Engineering & Lakehouse: Azure Databricks, Apache Spark, PySpark, Spark SQL, Delta Lake, Medallion Architecture, ETL Pipelines, ELT Pipelines, Data Lakehouse, ADLS Gen2, Azure Synapse Analytics, Microsoft Fabric, Change Data Capture (CDC), Unity Catalog

Data Modeling & Databases: SQL, PostgreSQL, Snowflake, Data Modeling, Star Schema, Redis, MongoDB, Query Optimization, Cloud Data Management

Streaming & Event Hubs: Apache Kafka, Spark Structured Streaming, Event Streaming, Kafka Connect, Azure Event Hubs

DataOps, IaC & CI/CD: Terraform, Infrastructure as Code (IaC), Azure DevOps, CI/CD, Docker, Kubernetes, Helm, Git, GitHub Actions, GitLab CI, Linux, Bash, dbt

Observability & Quality: Great Expectations, Data Quality Gates, Prometheus, Grafana, OpenTelemetry, Data Pipelines Observability, Python (Programming Language)

PROFESSIONAL EXPERIENCE

CGI India (Client: PNC Bank GCC) **Jun 2026 – Present**

DevOps Engineer / Azure Data Engineer *Hyderabad / Bangalore*

- **Engineered** scalable cloud data lakehouse infrastructure on Azure using Terraform modules, establishing automated ADLS Gen2 storage tiers, Azure Databricks workspaces, and Unity Catalog RBAC governance across 3 environments, cutting provisioning time by **60%**.
- **Orchestrated** robust batch ETL/ELT pipelines using Azure Data Factory (ADF) and PySpark with complex Spark SQL transformations, slashing daily financial dataset ingestion runtime by **40%** across 8 production data feeds.
- **Automated** declarative CI/CD release pipelines across Azure DevOps and GitLab CI with Trivy container vulnerability scanning and automated data validation quality gates, cutting deployment lead time by **45%**.

HFCL Limited **Aug 2025 – Jun 2026**

Engineer 1 — DevOps / Platform Engineering *Gurugram / Bangalore*

- **Orchestrated** stateful distributed data backbones (Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB) on Kubernetes via StatefulSets; configured Velero DR and Prometheus/Grafana alerts, slashing **RTO from 8h to 30min (93%)** and **MTTR by 45%**.
- **Automated** bare-metal Kubernetes (kubeadm) cluster lifecycle management using 40+ Ansible roles and Jinja2 templates (OS hardening, CNI, Harbor, Helm), cutting cluster bootstrap time **from 2 days to 4 hours (83% reduction)**.
- **Provisioned** and monitored a 50-node on-premise Kubernetes data platform for 15+ mission-critical microservices with namespace isolation, RBAC, and HPA, sustaining **99.95% platform availability**.

CGI India **Sep 2024 – Aug 2025**

Security Analyst Apprentice (Log Analytics & Telemetry) *Bangalore, India*

- **Engineered** centralized log telemetry ingestion pipelines via ELK Stack (Elasticsearch, Logstash, Kibana) and SQL analytical queries, authoring custom Logstash Grok parsing patterns that transformed **500GB+** of unstructured telemetry into structured schemas, accelerating incident triage by **50%**.

Stipe Solutions **Mar 2024 – Sep 2024**

Full Stack Developer Intern *Bengaluru, India*

- **Shipped** 30+ responsive UI screens and data validation workflows using React Native and Flutter, integrating RESTful APIs and SQL queries for **1,000+ daily active users**, cutting initial screen load latency by **35%**.
- **Optimized** client-side caching and mobile state management across iOS and Android builds, slashing application crash rates by **25%**.

KEY ENGINEERING PROJECTS

Enterprise Medallion Lakehouse Platform on Azure | *Azure Databricks, PySpark, Delta Lake, ADF, ADLS Gen2, SQL* **2026**

- **Architected** an end-to-end Medallion Lakehouse on Azure Databricks and Delta Lake, orchestrating ADF pipelines to ingest **1.5M+ events/day** into curated Bronze, Silver, and Gold layers across ADLS Gen2 with Unity Catalog governance.
- **Optimized** distributed PySpark and Spark SQL batch transformations with schema validation and partition pruning, slashing ETL pipeline runtime **from 4.0h to 2.6h (35% reduction)** while ensuring ACID consistency.

Real-Time Streaming Ingestion & CDC Lakehouse | *Apache Kafka, Spark Structured Streaming, Delta Lake, Unity Catalog, SQL* **2025**

- **Engineered** a low-latency real-time streaming ingestion pipeline using Apache Kafka and Spark Structured Streaming, processing **2M+ events/day** into Delta Lake with **under 500ms** end-to-end ingestion latency.
- **Integrated** Change Data Capture (CDC) with Delta Lake MERGE operations and Unity Catalog access controls, cutting data freshness lag **from 4 hours to sub-minute (98% reduction)**.

EDUCATION

B.Tech in Computer Science and Engineering | Maharishi Dayanand University (MDU), Rohtak (2020–2023)

CERTIFICATIONS

Microsoft Certified: Azure Fundamentals (AZ-900) | Microsoft SC-200 | Databricks Data Engineer Associate (In Progress)