Lavkush Kumar

Hyderabad, India | +91-9122036484 | lavkushry@gmail.com | linkedin.com/in/lavkushry | github.com/lavkushry

PROFESSIONAL SUMMARY

Azure Data Engineer with 2.5 YOE specializing in Azure Databricks, PySpark, Delta Lake, and ADF within enterprise Medallion Lakehouses. Engineered scalable batch and real-time streaming architectures (Apache Kafka, ADLS Gen2, Unity Catalog) processing 2M+ events/day into Bronze, Silver, and Gold Delta tables with sub-second latency. Automated DataOps CI/CD release cycles and infrastructure provisioning via Terraform, cutting ETL execution runtime by 35% and sustaining 99.95% cloud data platform availability.

TECHNICAL SKILLS

Cloud & Azure Data Services: Azure Data Factory (ADF), Azure Databricks, ADLS Gen2, Azure Synapse Analytics, Microsoft Fabric, AWS (EKS, S3, IAM)

Big Data & Lakehouse: PySpark, Apache Spark (Spark SQL), Delta Lake (ACID, Time Travel, Liquid Clustering), Medallion Architecture (Bronze/Silver/Gold), ETL/ELT Pipelines

Streaming & Databases: Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB, advanced SQL (window functions, CTEs), Star Schema Data Modeling

DataOps, IaC & CI/CD: Terraform (IaC), Docker, Kubernetes (kubeadm, EKS), Helm, Azure DevOps, GitLab CI, Jenkins, Great Expectations

Observability & Governance: Prometheus, Grafana, OpenTelemetry, ELK Stack (Logstash Grok), Unity Catalog RBAC, Linux/Bash, Python, Git

PROFESSIONAL EXPERIENCE

CGI India (Client: PNC Bank GCC) **Jun 2026 – Present**

DevOps Engineer / Azure Data Engineer *Hyderabad / Bangalore*

- **Engineered** scalable cloud data lakehouse infrastructure on Azure using Terraform modules, establishing automated ADLS Gen2 storage tiers, Azure Databricks workspaces, and Unity Catalog RBAC governance across 3 environments, cutting provisioning time by **60%**.
- **Orchestrated** robust batch ETL/ELT pipelines using Azure Data Factory (ADF) and PySpark with complex Spark SQL transformations, slashing daily financial dataset ingestion runtime by **40%** across 8 production data feeds.
- **Automated** declarative CI/CD release pipelines across Azure DevOps and GitLab CI with Trivy container vulnerability scanning and automated data validation quality gates, cutting deployment lead time by **45%**.

HFCL Limited **Aug 2025 – Jun 2026**

Engineer 1 — DevOps / Platform Engineering *Gurugram / Bangalore*

- **Orchestrated** stateful distributed data backbones (Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB) on Kubernetes via StatefulSets; configured Velero DR and Prometheus/Grafana alerts, slashing **RTO from 8h to 30min (93%)** and **MTTR by 45%**.
- **Automated** bare-metal Kubernetes (kubeadm) cluster lifecycle management using 40+ Ansible roles and Jinja2 templates (OS hardening, CNI, Harbor, Helm), cutting cluster bootstrap time **from 2 days to 4 hours (83% reduction)**.
- **Provisioned** and monitored a 50-node on-premise Kubernetes data platform for 15+ mission-critical microservices with namespace isolation, RBAC, and HPA, sustaining **99.95% platform availability**.

CGI India **Sep 2024 – Aug 2025**

Security Analyst Apprentice (Log Analytics & Telemetry) *Bangalore, India*

- **Engineered** centralized log telemetry ingestion pipelines via ELK Stack (Elasticsearch, Logstash, Kibana) and SQL analytical queries, authoring custom Logstash Grok parsing patterns that transformed **500GB+** of unstructured telemetry into structured schemas, accelerating incident triage by **50%**.

Stipe Solutions **Mar 2024 – Sep 2024**

Full Stack Developer Intern *Bengaluru, India*

- **Shipped** 30+ responsive UI screens and data validation workflows using React Native and Flutter, integrating RESTful APIs and SQL queries for **1,000+ daily active users**, cutting initial screen load latency by **35%**.
- **Optimized** client-side caching and mobile state management across iOS and Android builds, slashing application crash rates by **25%**.

KEY ENGINEERING PROJECTS

Enterprise Medallion Lakehouse Platform on Azure | *Azure Databricks, PySpark, Delta Lake, ADF, ADLS Gen2, SQL* **2026**

- **Architected** an end-to-end Medallion Lakehouse on Azure Databricks and Delta Lake, orchestrating ADF pipelines to ingest **1.5M+ events/day** into curated Bronze, Silver, and Gold layers across ADLS Gen2 with Unity Catalog governance.
- **Optimized** distributed PySpark and Spark SQL batch transformations with schema validation and partition pruning, slashing ETL pipeline runtime **from 4.0h to 2.6h (35% reduction)** while ensuring ACID consistency.

Real-Time Streaming Ingestion & CDC Lakehouse | *Apache Kafka, Spark Structured Streaming, Delta Lake, Unity Catalog, SQL* **2025**

- **Engineered** a low-latency real-time streaming ingestion pipeline using Apache Kafka and Spark Structured Streaming, processing **2M+ events/day** into Delta Lake with **under 500ms** end-to-end ingestion latency.
- **Integrated** Change Data Capture (CDC) with Delta Lake MERGE operations and Unity Catalog access controls, cutting data freshness lag **from 4 hours to sub-minute (98% reduction)**.

EDUCATION

B.Tech in Computer Science and Engineering | Maharishi Dayanand University (MDU), Rohtak (2020–2023)

CERTIFICATIONS

Microsoft Certified: Azure Fundamentals (AZ-900) | Microsoft SC-200 | Databricks Data Engineer Associate (In Progress)