LAVKUSH KUMAR

Hyderabad, India • +91 9122036484 • lavkushry@gmail.com • linkedin.com/in/lavkushry • github.com/lavkushry

TECHNICAL COMPETENCIES

Data Platforms & Lakehouse: Azure Databricks, PySpark (Spark SQL), Delta Lake (ACID, Time Travel, Liquid Clustering), Azure Data Factory (ADF), ADLS Gen2, Medallion Architecture (Bronze/Silver/Gold).
Streaming & Databases: Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB, advanced SQL (window functions, CTEs).
Cloud & DataOps: Terraform (IaC), AWS (EKS, VPC, IAM), Docker, Helm, Jenkins CI/CD, GitLab CI, Linux/Bash, Git.
Observability & Governance: Prometheus, Grafana, ELK Stack (Logstash Grok), Data Quality Gates, Unity Catalog concepts.

PROFESSIONAL EXPERIENCE

CGI India (PNC Bank GCC) | Azure Data Engineer / Cloud Infrastructure (Jun 2026 – Present • Hyderabad / Bangalore)

- Engineered scalable cloud data lakehouse infrastructure on Azure using Terraform modules, establishing automated ADLS Gen2 storage tiers, Azure Databricks workspaces, and Unity Catalog RBAC governance across 3 environments, cutting provisioning time by 60%.
- Orchestrated robust batch ETL/ELT pipelines using Azure Data Factory (ADF) and PySpark with complex Spark SQL transformations, slashing daily financial dataset ingestion runtime by 40% across 8 production data feeds.
- Automated declarative CI/CD release pipelines across Azure DevOps and GitLab CI with Trivy container vulnerability scanning and automated data validation quality gates, cutting deployment lead time by 45%.

Himachal Futuristic Communications (HFCL) | Engineer 1 - DevOps - Platform Engineering (Aug 2025 – Jun 2026 • Gurugram / Bangalore)

- Orchestrated stateful distributed data backbones (Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB) on Kubernetes via StatefulSets; configured Velero DR and Prometheus/Grafana alerts, slashing RTO from 8h to 30min (93% reduction) and MTTR by 45%.
- Automated bare-metal Kubernetes (kubeadm) cluster lifecycle management using 40+ Ansible roles and Jinja2 templates (OS hardening, CNI, Harbor, Helm), cutting cluster bootstrap time from 2 days to 4 hours (83% reduction).
- Provisioned and monitored a 50-node on-premise Kubernetes data platform for 15+ mission-critical microservices with namespace isolation, RBAC, and HPA, sustaining 99.95% platform availability.

CGI India | Security Analyst Apprentice (Log Analytics & Telemetry) (Sep 2024 – Aug 2025 • Bangalore)

- Engineered centralized log telemetry ingestion pipelines via ELK Stack (Elasticsearch, Logstash, Kibana) and SQL analytical queries, authoring custom Logstash Grok parsing patterns that transformed 500GB+ of unstructured telemetry into structured schemas, accelerating incident triage by 50%.

Stipe Solution Private Limited | Full Stack Developer Intern (Mar 2024 – Sep 2024 • Bengaluru)

- Shipped 30+ responsive UI screens and data validation workflows using React Native and Flutter, integrating RESTful APIs and SQL queries for 1,000+ daily active users, cutting initial screen load latency by 35%.
- Optimized client-side caching and mobile state management across iOS and Android builds, slashing application crash rates by 25%.

KEY ENGINEERING PROJECTS

Enterprise Medallion Lakehouse Platform on Azure (Azure Databricks, PySpark, Delta Lake, ADF, ADLS Gen2, Terraform, SQL)

- Architected an end-to-end Medallion Lakehouse on Azure Databricks and Delta Lake, orchestrating ADF pipelines to ingest 1.5M+ events/day into curated Bronze, Silver, and Gold layers across ADLS Gen2 with Unity Catalog governance.
- Optimized distributed PySpark and Spark SQL batch transformations with schema validation and partition pruning, slashing ETL pipeline runtime from 4.0h to 2.6h (35% reduction) while ensuring ACID transactional consistency.

Real-Time Streaming Ingestion & CDC Lakehouse (Apache Kafka, Spark Structured Streaming, Delta Lake, Unity Catalog, SQL)

- Engineered a low-latency real-time streaming ingestion pipeline using Apache Kafka and Spark Structured Streaming, processing 2M+ events/day into Delta Lake with under 500ms end-to-end ingestion latency.
- Integrated Change Data Capture (CDC) with Delta Lake MERGE operations and Unity Catalog access controls, cutting data freshness lag from 4 hours to sub-minute latency (98% reduction).

Cloud Data Platform CI/CD & Governance (Azure DevOps, Terraform, Great Expectations, Prometheus/Grafana, Databricks)

- Automated end-to-end DataOps release cycles using Azure DevOps multi-stage CI/CD pipelines and Terraform IaC across 3 environments, cutting workspace deployment time from 6 hours to 20 minutes (94% reduction).
- Configured automated data quality validation gates with Great Expectations and Prometheus/Grafana monitoring, catching 99.2% of upstream schema anomalies before Gold table ingestion.

Cloud AI Infrastructure & Observability Platform (Kubernetes, Docker, Prometheus, Grafana, OpenTelemetry, Python, CI/CD)

- Built an end-to-end telemetry and observability backbone using Prometheus, Grafana, and OpenTelemetry, monitoring 15+ distributed microservices with under 100ms scraping latency, slashing incident MTTR by 40%.
- Automated deployment and scaling of data serving microservices using Docker, Kubernetes, and GitHub Actions with automated canary rollback safeguards, sustaining 99.9% platform availability.

EDUCATION & CERTIFICATIONS

B.Tech in Computer Science and Engineering | Maharishi Dayanand University (MDU), Rohtak (2020 – 2023)
Certifications: Microsoft Certified: Azure Fundamentals (AZ-900) | Microsoft SC-200 | Databricks Data Engineer Associate (In Progress)