# LAVKUSH KUMAR

• +91 9122036484 • lavkushry@gmail.com • LinkedIn • GitHub

## TECHNICAL COMPETENCIES

**Data Platforms & Lakehouse:** A, z, u, r, e, , D, a, t, a, b, r, i, c, k, s, ,, , P, y, S, p, a, r, k, , (, S, p, a, r, k, , S, Q, L, ), ,, , D, e, l, t, a, , L, a, k, e, , (, A, C, I, D, ,, , T, i, m, e, , T, r, a, v, e, l, ,, , L, i, q, u, i, d, , C, l, u, s, t, e, r, i, n, g, ), ,, , A, z, u, r, e, , D, a, t, a, , F, a, c, t, o, r, y, , (, A, D, F, ), ,, , A, D, L, S, , G, e, n, 2, ,, , M, e, d, a, l, l, i, o, n, , A, r, c, h, i, t, e, c, t, u, r, e, , (, B, r, o, n, z, e, /, S, i, l, v, e, r, /, G, o, l, d, ), .
**Streaming & Databases:** A, p, a, c, h, e, , K, a, f, k, a, ,, , R, e, d, i, s, ,, , P, o, s, t, g, r, e, S, Q, L, ,, , T, i, m, e, s, c, a, l, e, D, B, ,, , M, o, n, g, o, D, B, ,, , a, d, v, a, n, c, e, d, , S, Q, L, , (, w, i, n, d, o, w, , f, u, n, c, t, i, o, n, s, ,, , C, T, E, s, ), .
**Cloud & DataOps:** T, e, r, r, a, f, o, r, m, , (, I, a, C, ), ,, , A, W, S, , (, E, K, S, ,, , V, P, C, ,, , I, A, M, ), ,, , D, o, c, k, e, r, ,, , H, e, l, m, ,, , J, e, n, k, i, n, s, , C, I, /, C, D, ,, , G, i, t, L, a, b, , C, I, ,, , L, i, n, u, x, /, B, a, s, h, ,, , G, i, t, .
**Observability & Governance:** P, r, o, m, e, t, h, e, u, s, ,, , G, r, a, f, a, n, a, ,, , E, L, K, , S, t, a, c, k, , (, L, o, g, s, t, a, s, h, , G, r, o, k, ), ,, , D, a, t, a, , Q, u, a, l, i, t, y, , G, a, t, e, s, ,, , U, n, i, t, y, , C, a, t, a, l, o, g, , c, o, n, c, e, p, t, s, .

## PROFESSIONAL EXPERIENCE

### CGI India (PNC Bank GCC) | Azure Data Engineer / Cloud Infrastructure ( • Hyderabad / Bangalore)

- Standardized multi-tenant AWS EKS cluster infrastructure using Terraform modules and Ansible, establishing reproducible VPC peering, IAM role bindings, and auto-scaling node groups across environments.
- Engineered declarative CI/CD pipelines across Jenkins and GitLab CI with Trivy vulnerability scanning and quality gates, cutting deployment lead time by 40%.
- Hardened ingress routing and data API gateway policies using Nginx Ingress and Kong Gateway with rate-limiting and mTLS, slashing gateway latency overhead by 15% across 8 production microservices.

### Himachal Futuristic Communications (HFCL) | Engineer 1 - DevOps - Platform Engineering ( • Gurugram / Bangalore)

- Orchestrated stateful distributed data backbones (Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB) on Kubernetes via StatefulSets; configured Velero DR and Prometheus/Grafana alerts, slashing RTO from 8h to 30min (93% reduction) and MTTR by 45%.
- Automated bare-metal Kubernetes (kubeadm) cluster lifecycle management using 40+ Ansible roles and Jinja2 templates (OS hardening, CNI, Harbor, Helm), cutting cluster bootstrap time from 2 days to 4 hours (83% reduction).
- Operated a 50-node on-premise Kubernetes platform for 15+ mission-critical microservices with namespace isolation, RBAC, and HPA, sustaining 99.95% platform availability.

### CGI India | Security Analyst Apprentice (Log Analytics & Telemetry) ( • Bangalore)

- Deployed centralized log telemetry pipelines via ELK Stack (Elasticsearch, Logstash, Kibana), authoring custom Logstash Grok patterns that transformed unstructured telemetry into structured schemas and accelerated incident triage by 50%.

## KEY ENGINEERING PROJECTS

### Enterprise Medallion Lakehouse Platform on Azure

- Architected an end-to-end Medallion Lakehouse on Azure Databricks and Delta Lake, orchestrating ADF pipelines to ingest transactional event datasets into curated Bronze, Silver, and Gold layers across ADLS Gen2.
- Optimized distributed PySpark batch transformations with schema validation and partition pruning, slashing ETL pipeline runtime by 35% while ensuring ACID transactional consistency.

### Real-Time Streaming Ingestion & CDC Lakehouse

- Engineered a low-latency real-time streaming ingestion pipeline using Apache Kafka and Spark Structured Streaming, processing 2M+ events/day into Delta Lake with under 500ms end-to-end ingestion latency.
- Implemented Change Data Capture (CDC) with Delta Lake merge operations and Unity Catalog access controls, cutting data freshness lag from 4 hours to near real-time (sub-minute).

### Cloud Data Platform CI/CD & Governance

- Automated end-to-end DataOps release cycles using Azure DevOps multi-stage YAML pipelines and Terraform IaC, provisioning reproducible Databricks workspaces and storage accounts with zero manual intervention.
- Configured automated data quality validation gates with Great Expectations and Prometheus/Grafana monitoring, catching 99.2% of upstream schema anomalies before Gold table ingestion.

## EDUCATION & CERTIFICATIONS

**B.Tech in Computer Science and Engineering** | Maharishi Dayanand University (MDU), Rohtak (2020 – 2023)
**Certifications:** Microsoft Certified: Azure Fundamentals (AZ-900) | Microsoft SC-200 | Databricks Data Engineer Associate (In Progress)