LAVKUSH KUMAR

Hyderabad, India • +91-9122036484 • lav84095@gmail.com • linkedin.com/in/lavkush-kumar • github.com/lavkush-kumar

PROFESSIONAL SUMMARY

Azure Data Engineer with 2.5 YOE specializing in Azure Databricks, PySpark, Delta Lake, and ADF within enterprise Medallion Lakehouses. Engineered scalable batch and real-time streaming architectures (Apache Kafka, ADLS Gen2, Unity Catalog) processing 2M+ events/day into Bronze, Silver, and Gold Delta tables with sub-second latency. Automated DataOps CI/CD release cycles and infrastructure provisioning via Terraform, cutting ETL execution runtime by 35% and sustaining 99.95% cloud data platform availability.

TECHNICAL COMPETENCIES

Data Engineering & Lakehouse: Azure Databricks, Apache Spark, PySpark, Spark SQL, Delta Lake, Medallion Architecture, ETL Pipelines, ELT Pipelines, Data Lakehouse, ADLS Gen2, Azure Synapse Analytics, Microsoft Fabric, Change Data Capture (CDC), Unity Catalog.
Data Modeling & Databases: SQL, PostgreSQL, Snowflake, Data Modeling, Star Schema, Redis, MongoDB, Query Optimization, Cloud Data Management.
Streaming & Event Hubs: Apache Kafka, Spark Structured Streaming, Event Streaming, Kafka Connect, Azure Event Hubs.
DataOps, IaC & CI/CD: Terraform, Infrastructure as Code (IaC), Azure DevOps, CI/CD, Docker, Kubernetes, Helm, Git, GitHub Actions, GitLab CI, Linux, Bash, dbt.
Observability & Data Quality: Great Expectations, Data Quality Gates, Prometheus, Grafana, OpenTelemetry, Data Pipelines Observability, Python (Programming Language).

PROFESSIONAL EXPERIENCE

CGI India (PNC Bank GCC) | DevOps Engineer / Azure Data Engineer **(Jun 2026 – Present • Hyderabad / Bangalore)**

- Engineered scalable cloud data lakehouse infrastructure on Azure using Terraform modules, establishing automated ADLS Gen2 storage tiers, Azure Databricks workspaces, and Unity Catalog RBAC governance across 3 environments, cutting provisioning time by 60%.
- Orchestrated robust batch ETL/ELT pipelines using Azure Data Factory (ADF) and PySpark with complex Spark SQL transformations, slashing daily financial dataset ingestion runtime by 40% across 8 production data feeds.
- Automated declarative CI/CD release pipelines across Azure DevOps and GitLab CI with Trivy container vulnerability scanning and automated data validation quality gates, cutting deployment lead time by 45%.

Himachal Futuristic Communications (HFCL) | Engineer 1 - DevOps - Platform Engineering **(Aug 2025 – Jun 2026 • Gurugram / Bangalore)**

- Orchestrated stateful distributed data backbones (Apache Kafka, Redis, PostgreSQL, TimescaleDB, MongoDB) on Kubernetes via StatefulSets; configured Velero DR and Prometheus/Grafana alerts, slashing RTO from 8h to 30min (93% reduction) and MTTR by 45%.
- Automated bare-metal Kubernetes (kubeadm) cluster lifecycle management using 40+ Ansible roles and Jinja2 templates (OS hardening, CNI, Harbor, Helm), cutting cluster bootstrap time from 2 days to 4 hours (83% reduction).
- Provisioned and monitored a 50-node on-premise Kubernetes data platform for 15+ mission-critical microservices with namespace isolation, RBAC, and HPA, sustaining 99.95% platform availability.

CGI India | Security Analyst Apprentice (Log Analytics & Telemetry) **(Sep 2024 – Aug 2025 • Bangalore)**

- Engineered centralized log telemetry ingestion pipelines via ELK Stack (Elasticsearch, Logstash, Kibana) and SQL analytical queries, authoring custom Logstash Grok parsing patterns that transformed 500GB+ of unstructured telemetry into structured schemas, accelerating incident triage by 50%.

Stipe Solution Private Limited | Full Stack Developer Intern **(Mar 2024 – Sep 2024 • Bengaluru)**

- Shipped 30+ responsive UI screens and data validation workflows using React Native and Flutter, integrating RESTful APIs and SQL queries for 1,000+ daily active users, cutting initial screen load latency by 35%.
- Optimized client-side caching and mobile state management across iOS and Android builds, slashing application crash rates by 25%.

KEY ENGINEERING PROJECTS

Enterprise Medallion Lakehouse Platform on Azure **(Azure Databricks, PySpark, Delta Lake, ADF, ADLS Gen2, Terraform, SQL)**

- Architected an end-to-end Medallion Lakehouse on Azure Databricks and Delta Lake, orchestrating ADF pipelines to ingest 1.5M+ events/day into curated Bronze, Silver, and Gold layers across ADLS Gen2 with Unity Catalog governance.
- Optimized distributed PySpark and Spark SQL batch transformations with schema validation and partition pruning, slashing ETL pipeline runtime from 4.0h to 2.6h (35% reduction) while ensuring ACID transactional consistency.

Real-Time Streaming Ingestion & CDC Lakehouse **(Apache Kafka, Spark Structured Streaming, Delta Lake, Unity Catalog, SQL)**

- Engineered a low-latency real-time streaming ingestion pipeline using Apache Kafka and Spark Structured Streaming, processing 2M+ events/day into Delta Lake with under 500ms end-to-end ingestion latency.
- Integrated Change Data Capture (CDC) with Delta Lake MERGE operations and Unity Catalog access controls, cutting data freshness lag from 4 hours to sub-minute latency (98% reduction).

Cloud Data Platform CI/CD & Governance **(Azure DevOps, Terraform, Great Expectations, Prometheus/Grafana, Databricks)**

- Automated end-to-end DataOps release cycles using Azure DevOps multi-stage CI/CD pipelines and Terraform IaC across 3 environments, cutting workspace deployment time from 6 hours to 20 minutes (94% reduction).
- Configured automated data quality validation gates with Great Expectations and Prometheus/Grafana monitoring, catching 99.2% of upstream schema anomalies before Gold table ingestion.

Cloud AI Infrastructure & Observability Platform **(Kubernetes, Docker, Prometheus, Grafana, OpenTelemetry, Python, CI/CD)**

- Built an end-to-end telemetry and observability backbone using Prometheus, Grafana, and OpenTelemetry, monitoring 15+ distributed microservices with under 100ms scraping latency, slashing incident MTTR by 40%.
- Automated deployment and scaling of data serving microservices using Docker, Kubernetes, and GitHub Actions with automated canary rollback safeguards, sustaining 99.9% platform availability.

EDUCATION

B.Tech in Computer Science and Engineering | Maharishi Dayanand University (MDU), Rohtak (2020 – 2023)

CERTIFICATIONS

Microsoft Certified: Azure Fundamentals (AZ-900) • Microsoft SC-200 • Databricks Data Engineer Associate (In Progress)